Convolution

Multi-channels input and multi-filters layer

Dr. Thanh-Sach LE
LTSACH@hcmut.edu.vn

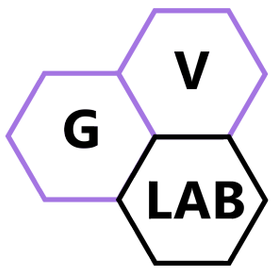

GVLab:
Graphics and Vision Laboratory

Faculty of Computer Science and Engineering,
HCMUT

2 Contents

- Multi-channels input
- Multi-filters layer

Convolution

Multi-channels input

Dr. Thanh-Sach LE
LTSACH@hcmut.edu.vn

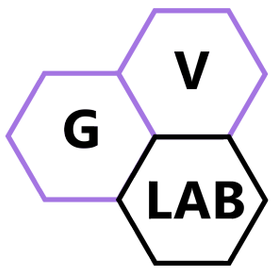

GVLab:
Graphics and Vision Laboratory

Faculty of Computer Science and Engineering,
HCMUT

Convolution algorithm

Start

- ✳ Rotate the kernel 180^0
- ✳ Flatten it to a vector

1

- ✳ Padding zero to input
- ✳ Allocate output buffer

2

- ✳ Align the kernel window with the top-left of the input

3

- ✳ Extract the sub-image overlapped with the kernel
- ✳ Flatten the sub-image to a vector
- ✳ Take dot-product with the flattened kernel
- ✳ Fill the result to output

4

all output values assigned?

No

- ✳ Slide the kernel window to next position on the input

5

Yes

End

5 Multi-channels input

Input image or feature map: multiple channels

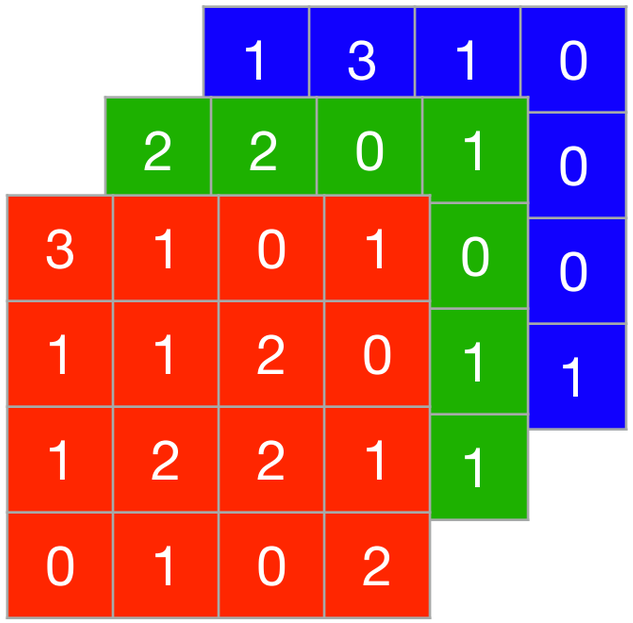

Input: 3 channels

6 Multi-channels input

Input image or feature map: multiple channels

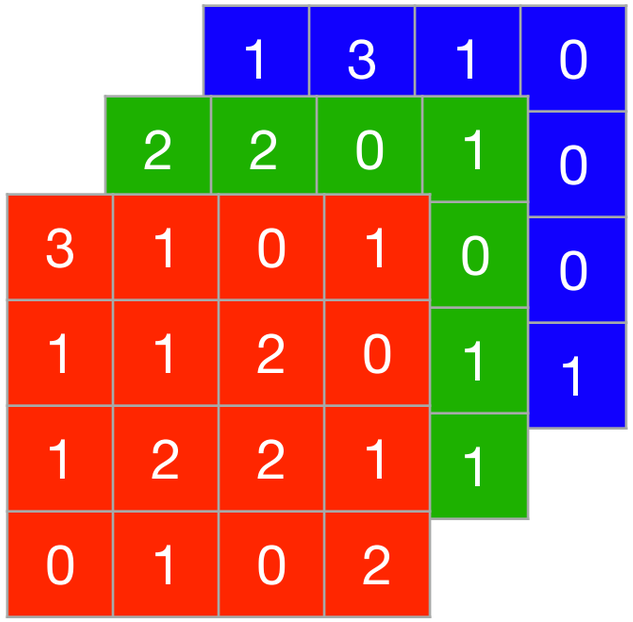

Input: 3 channels

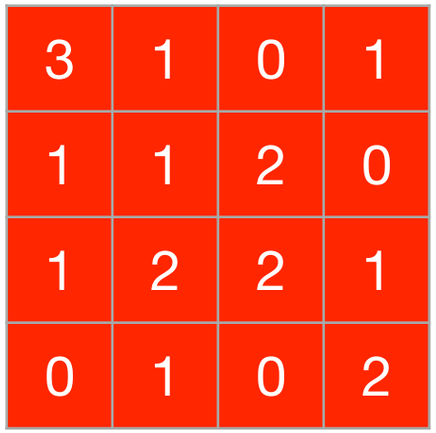

Channel 1 (RED)

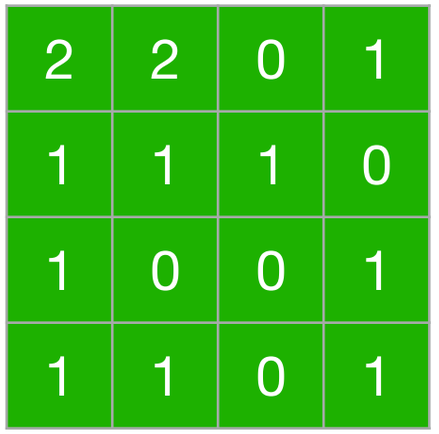

Channel 2 (GREEN)

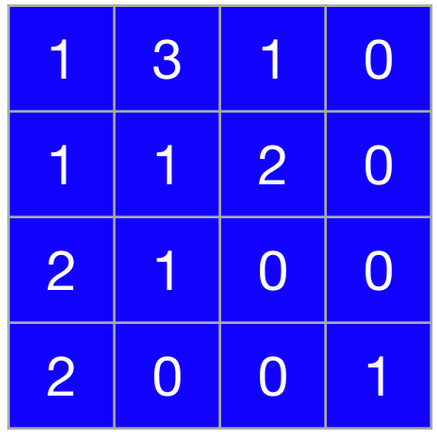

Channel 3 (BLUE)

7 Multi-channels input

Input image or feature map: multiple channels

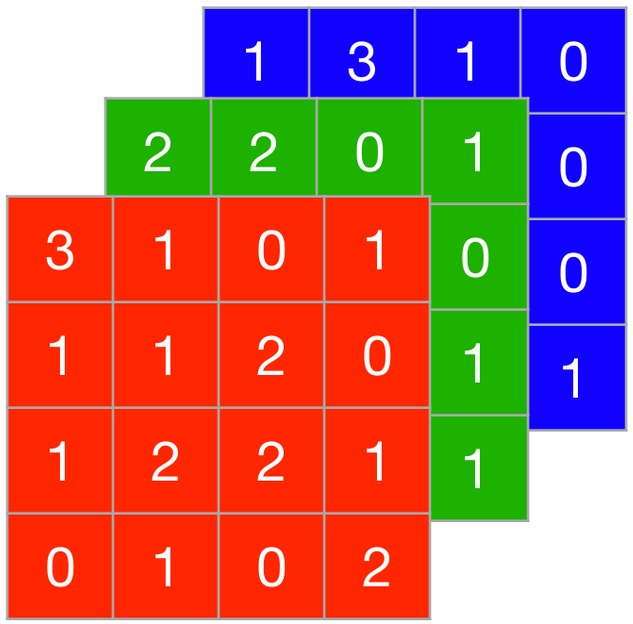

Input: 3 channels

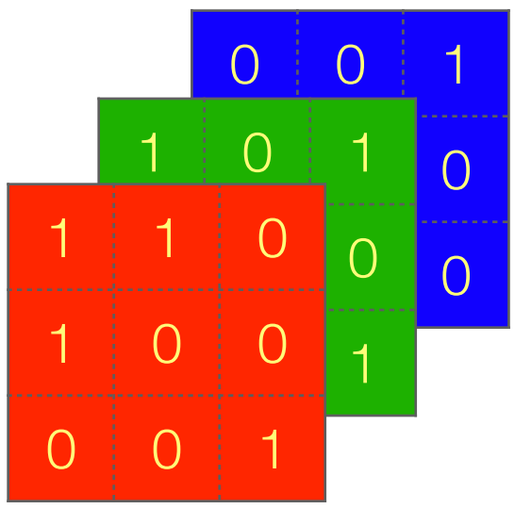

Kernel: 3 channels

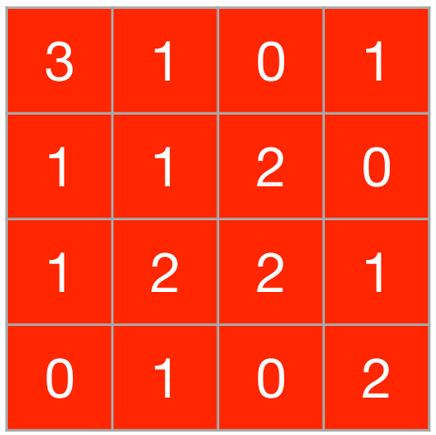

Channel 1 (RED)

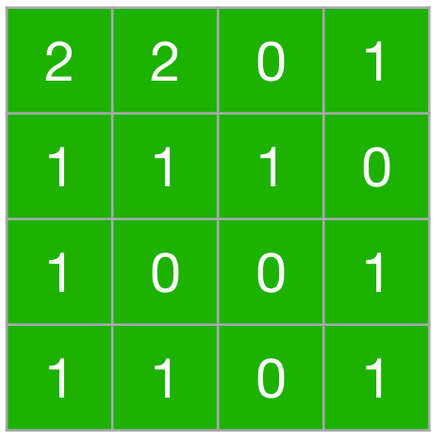

Channel 2 (GREEN)

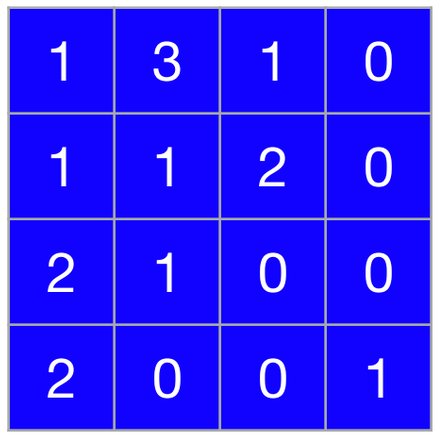

Channel 3 (BLUE)

8 Multi-channels input

Input image or feature map: multiple channels

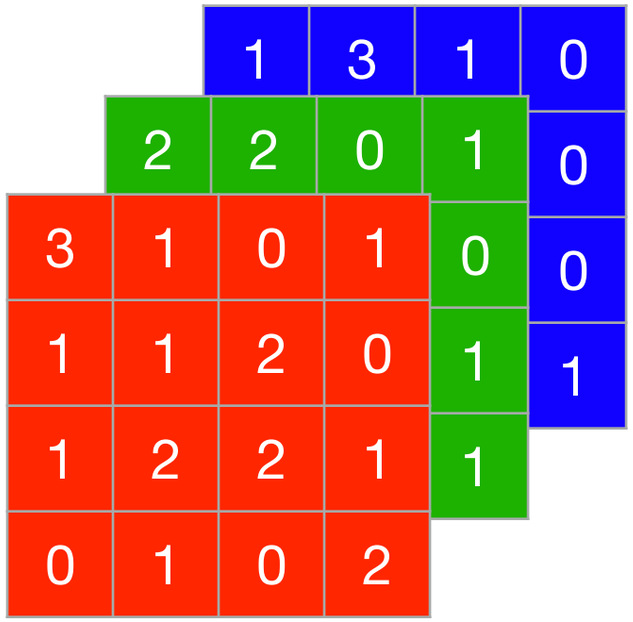

Input: 3 channels

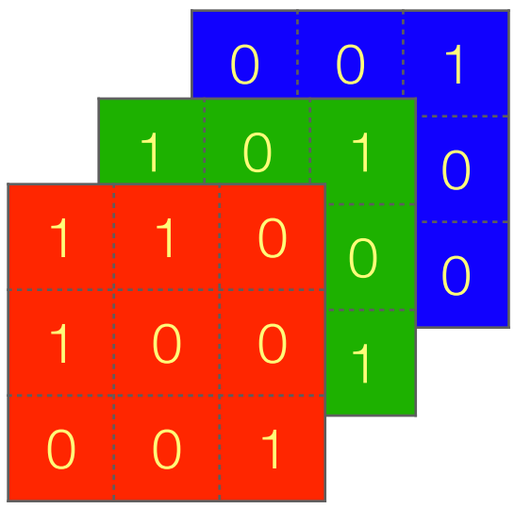

Kernel: 3 channels

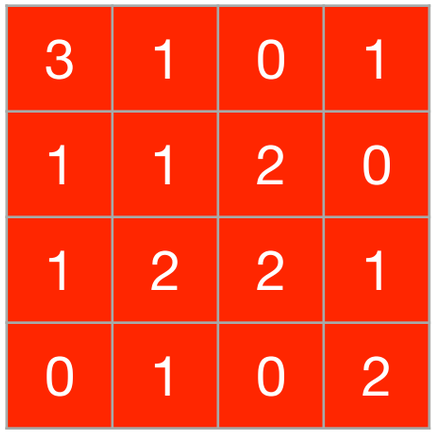

Channel 1 (RED)

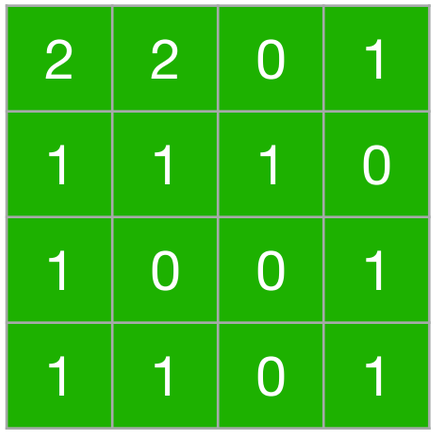

Channel 2 (GREEN)

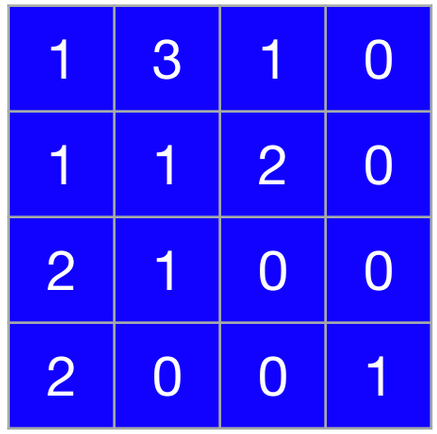

Channel 3 (BLUE)

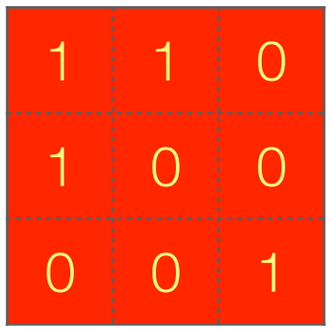

Channel 1 (RED)

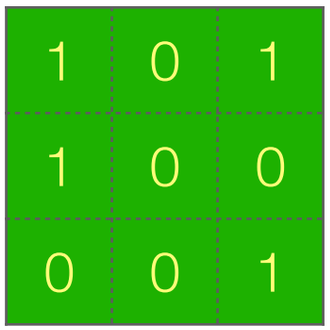

Channel 2 (GREEN)

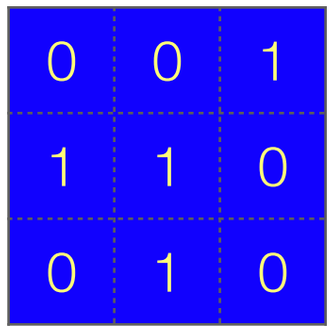

Channel 3 (BLUE)

9 Multi-channels input

1 Rotate kernel 180⁰

1	1	0
1	0	0
0	0	1

Rotate 180⁰

1	0	0
0	0	1
0	1	1

1	0	1
1	0	0
0	0	1

Rotate 180⁰

1	0	0
0	0	1
1	0	1

0	0	1
1	1	0
0	1	0

Rotate 180⁰

0	1	0
0	1	1
1	0	0

10 Multi-channels input

1 Flatten the rotated kernel to a vector

1	1	0
1	0	0
0	0	1

Rotate 180⁰

1	0	0
0	0	1
0	1	1

1	0	1
1	0	0
0	0	1

Rotate 180⁰

1	0	0
0	0	1
1	0	1

0	0	1
1	1	0
0	1	0

Rotate 180⁰

0	1	0
0	1	1
1	0	0

After flattening:

1	0	0	0	0	1	0	1	1	1	0	0	0	0	1	1	0	1	0	1	0	0	1	1	1	0	0

11 Multi-channels input

2 Padding the input

3	1	0	1
1	1	2	0
1	2	2	1
0	1	0	2

Channel 1 (RED)

2	2	0	1
1	1	1	0
1	0	0	1
1	1	0	1

Channel 2 (GREEN)

1	3	1	0
1	1	2	0
2	1	0	0
2	0	0	1

Channel 3 (BLUE)

(No padding)

12 Multi-channels input

2 Allocate the output buffer

3	1	0	1
1	1	2	0
1	2	2	1
0	1	0	2

Channel 1 (RED)

2	2	0	1
1	1	1	0
1	0	0	1
1	1	0	1

Channel 2 (GREEN)

1	3	1	0
1	1	2	0
2	1	0	0
2	0	0	1

Channel 3 (BLUE)

input: $i_1 = i_2 = 4$

kernel: $k_1 = k_2 = 3$

padding: $p_1 = p_2 = 0$

strides: $s_1 = s_2 = 1$

$i_1 - k_1 + 1 = 2$

$i_2 - k_2 + 1 = 2$

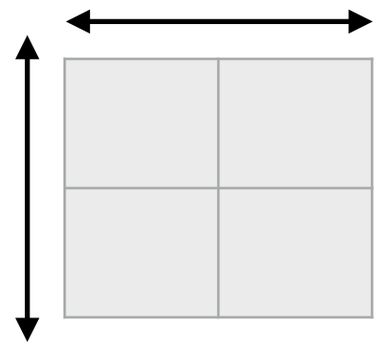

Output

3	1	0	1
1	1	2	0
1	2	2	1
0	1	0	2

2	2	0	1
1	1	1	0
1	0	0	1
1	1	0	1

1	3	1	0
1	1	2	0
2	1	0	0
2	0	0	1

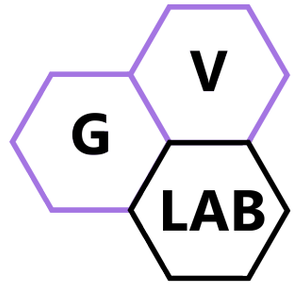

3 Starting the cross-correlation process

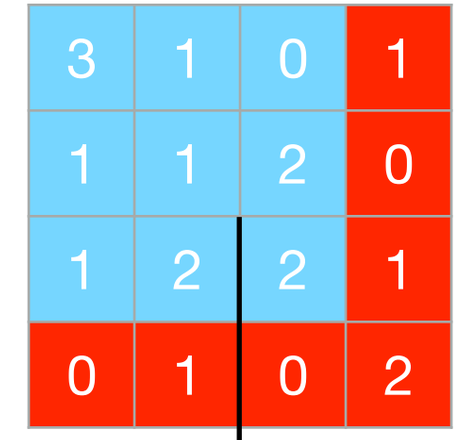

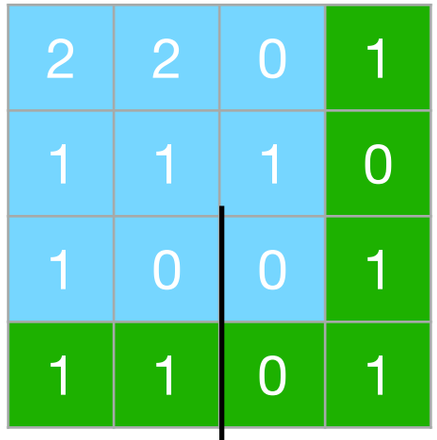

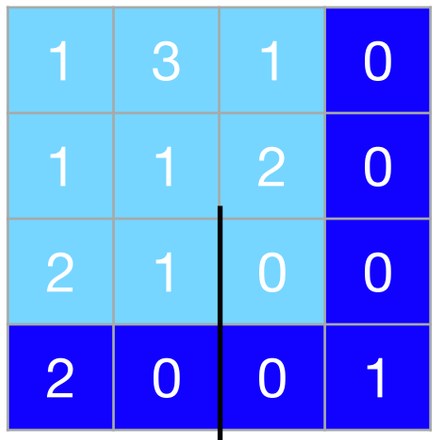

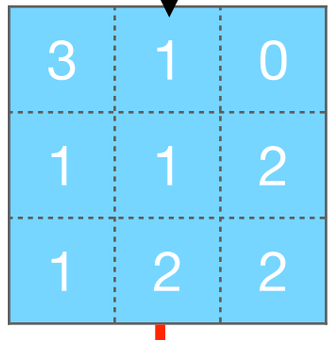

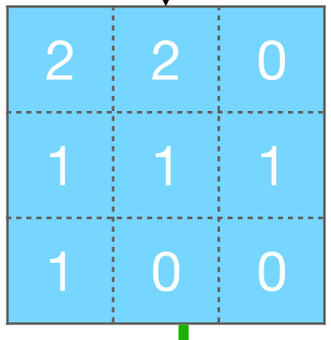

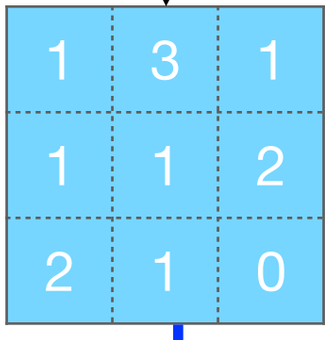

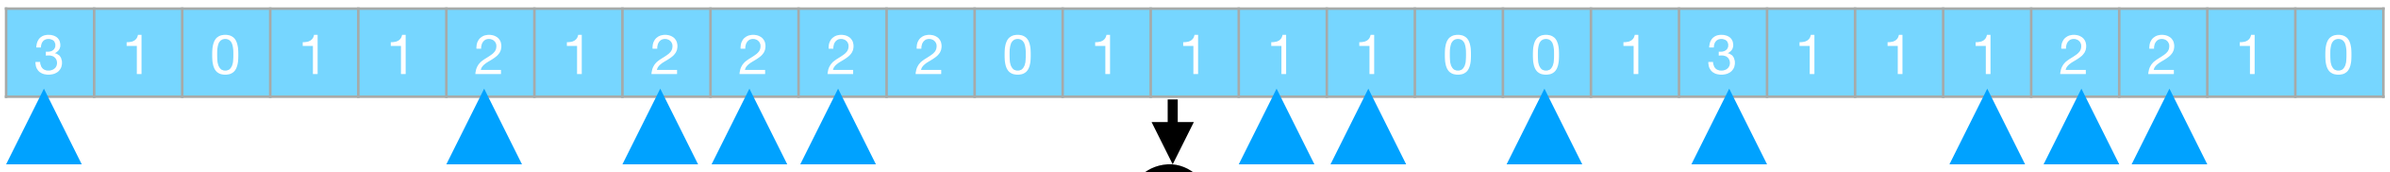

dot-product

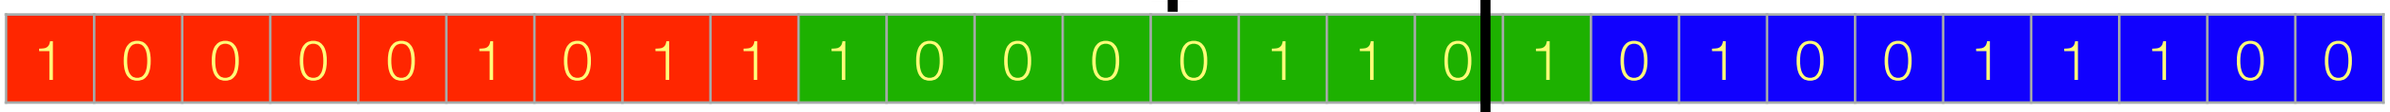

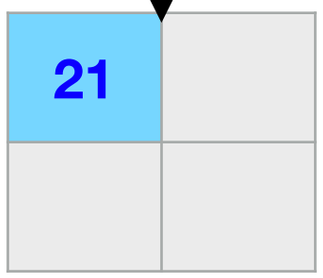

> Slide the kernel to right

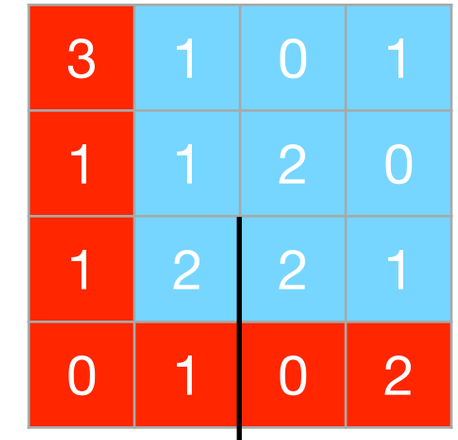

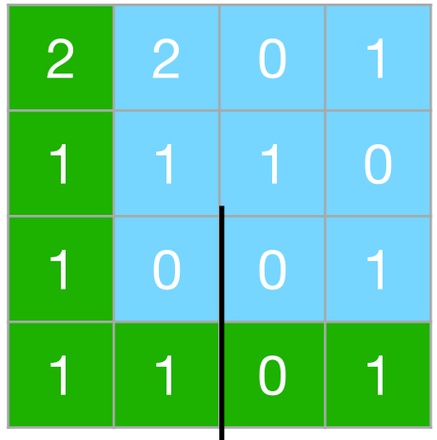

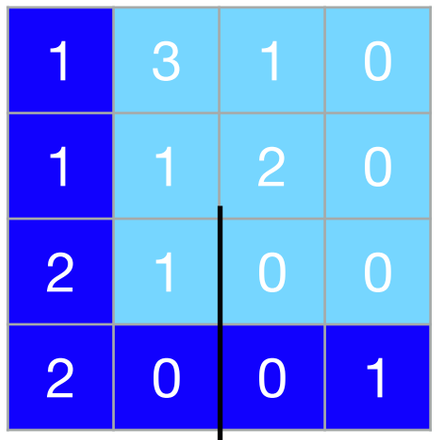

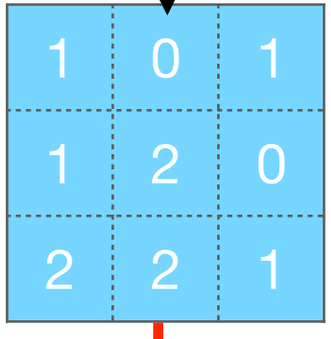

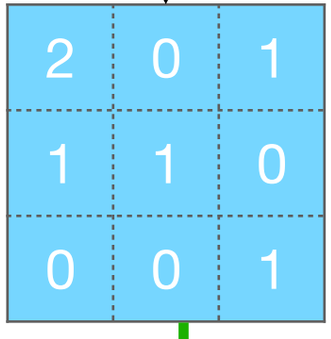

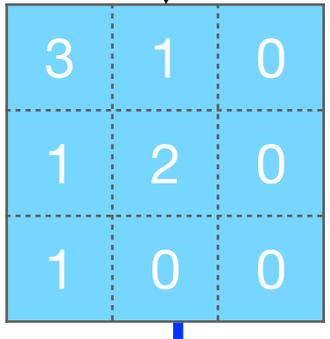

dot-product

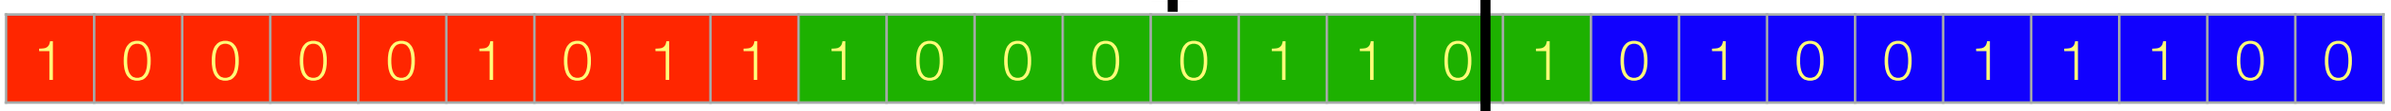

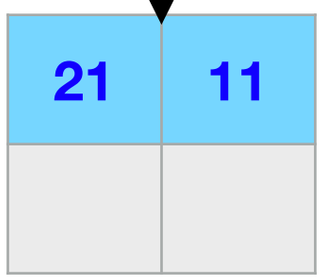

> Slide the kernel down

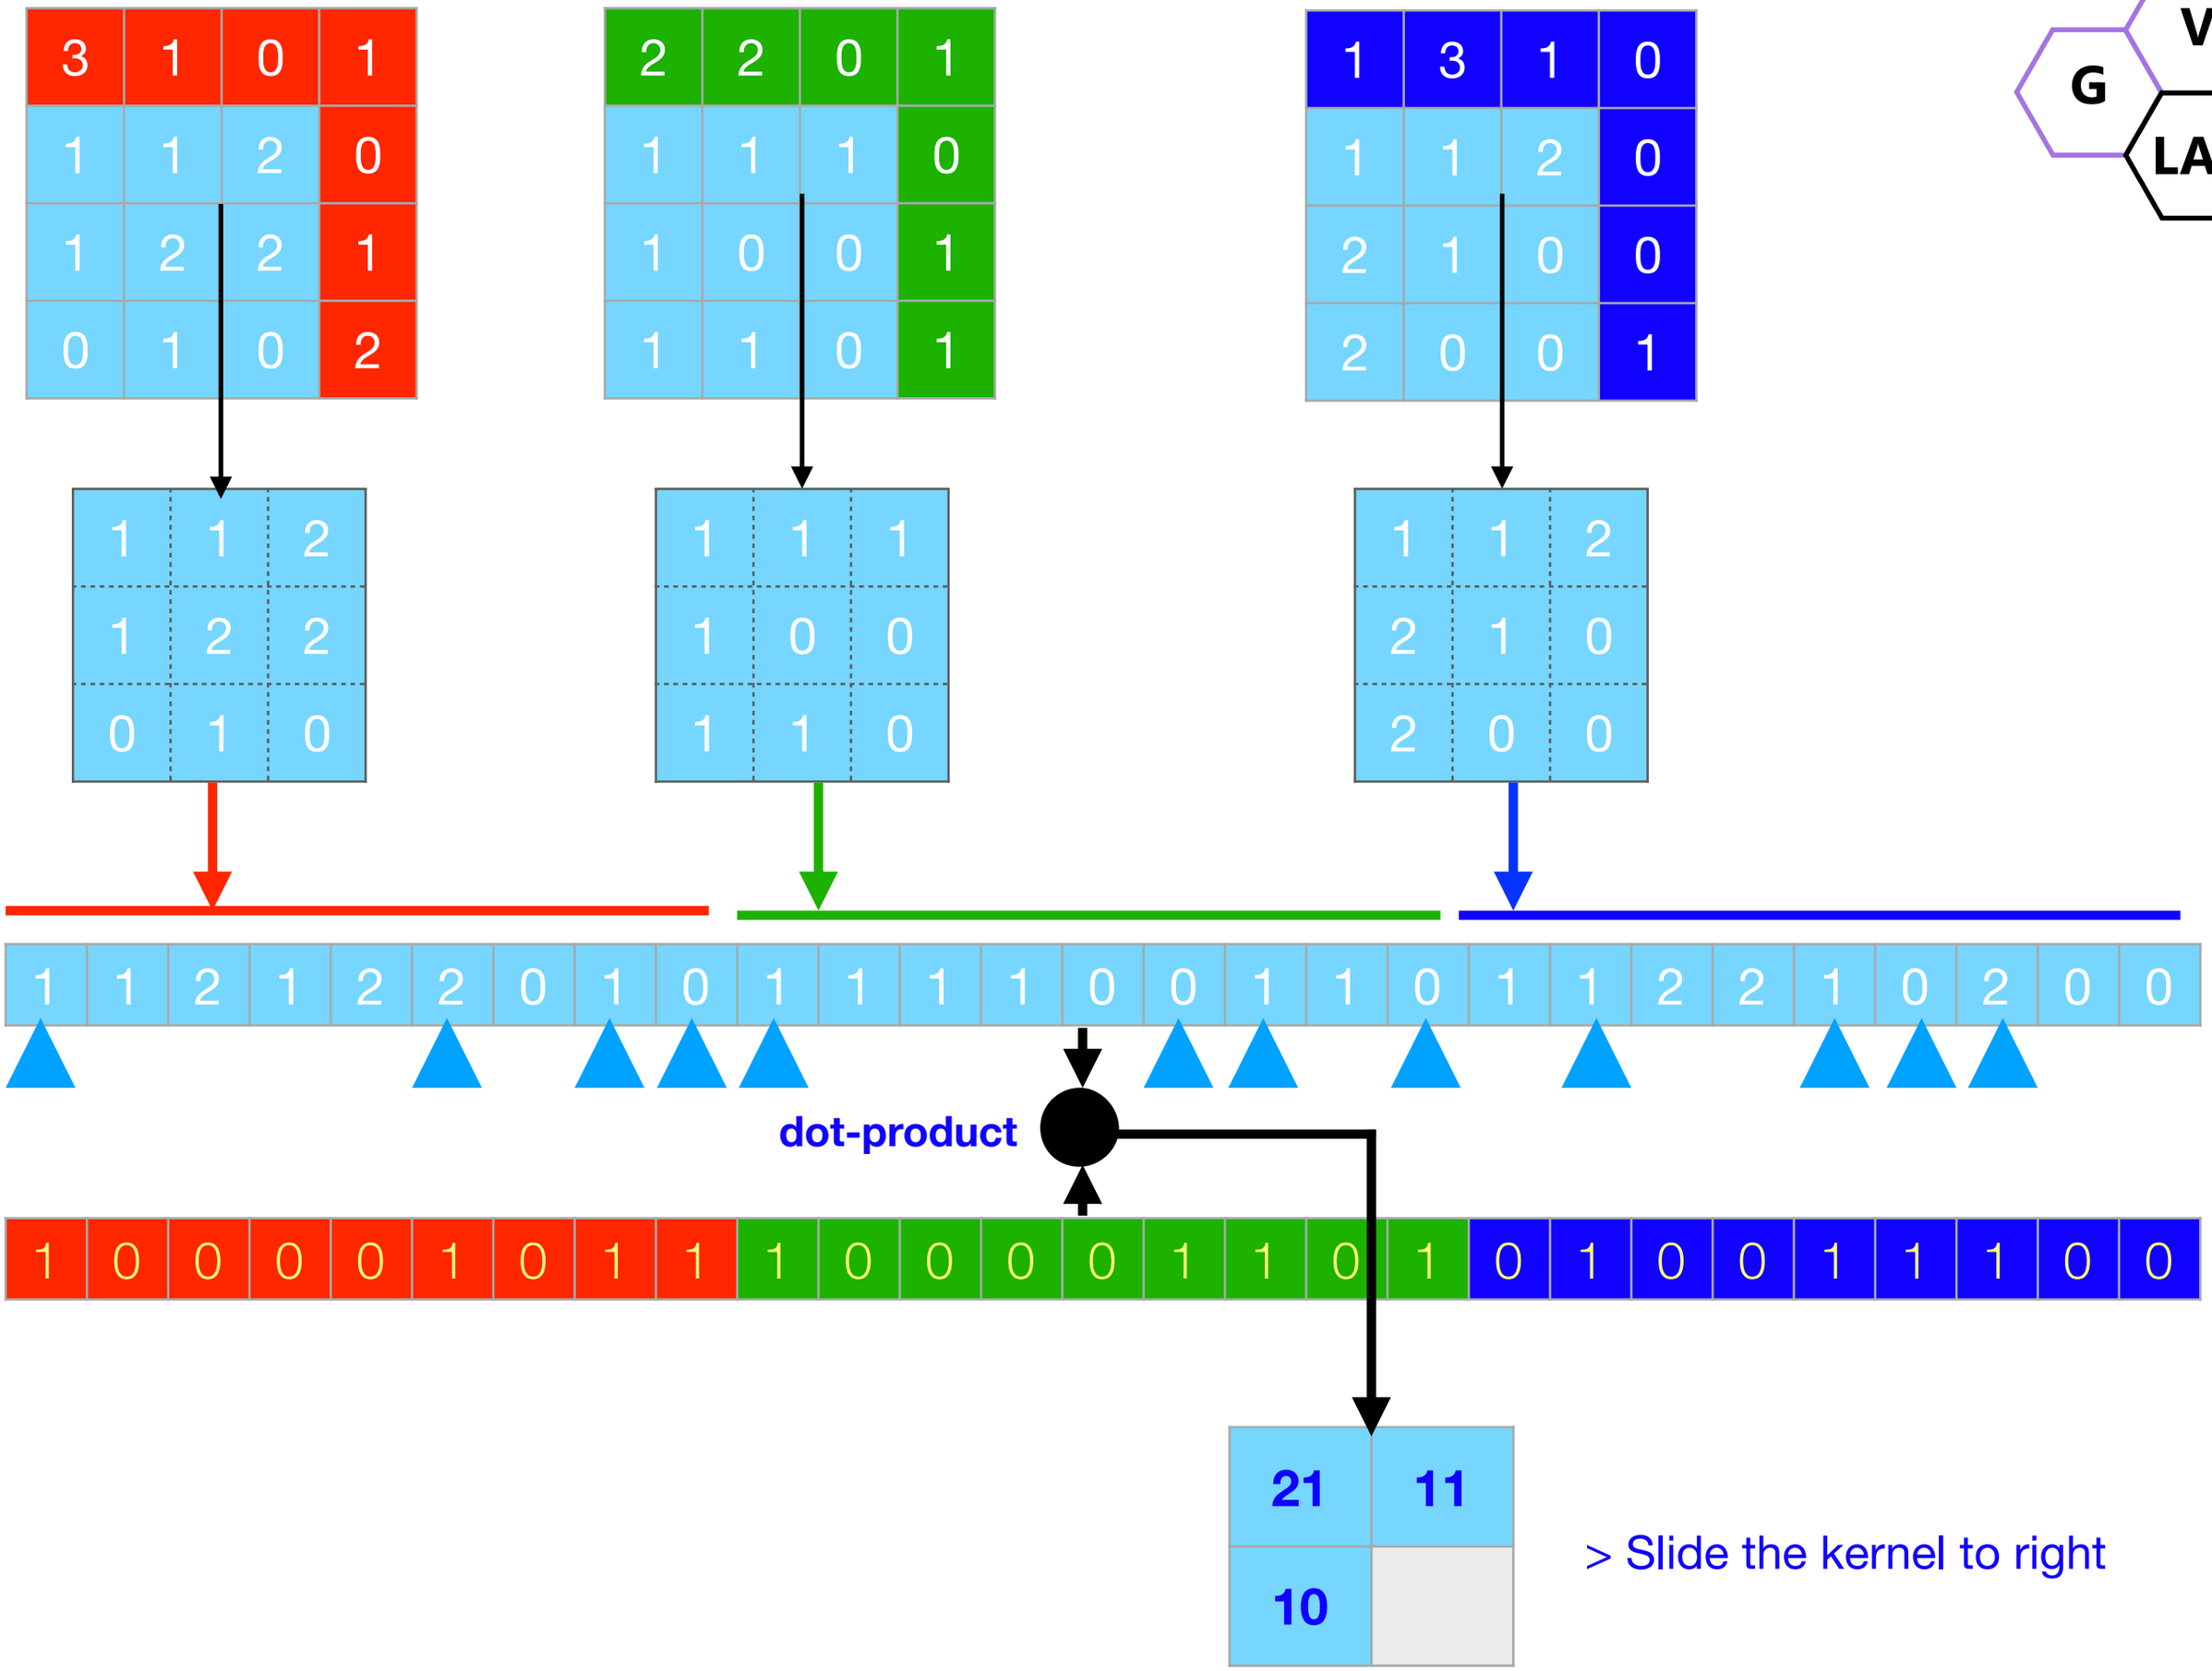
3 1 0 1
1 1 2 0
1 2 2 1
0 1 0 2
2 2 0 1
1 1 1 0
1 0 0 1
1 1 0 1
1 3 1 0
1 1 2 0
2 1 0 0
2 0 0 1
1 1 2
1 2 2
0 1 0
1 1 1
1 0 0
1 1 0
1 1 2
2 1 0
2 0 0
1 1 2 1 2 2 0 1 0 1 1 1 1 0 0 1 1 0 1 1 2 2 1 0 2 0 0
dot-product
1 0 0 0 0 1 0 1 1 1 0 0 0 0 1 1 0 1 0 1 0 0 1 1 1 0 0
21
11
10
> Slide the kernel to right

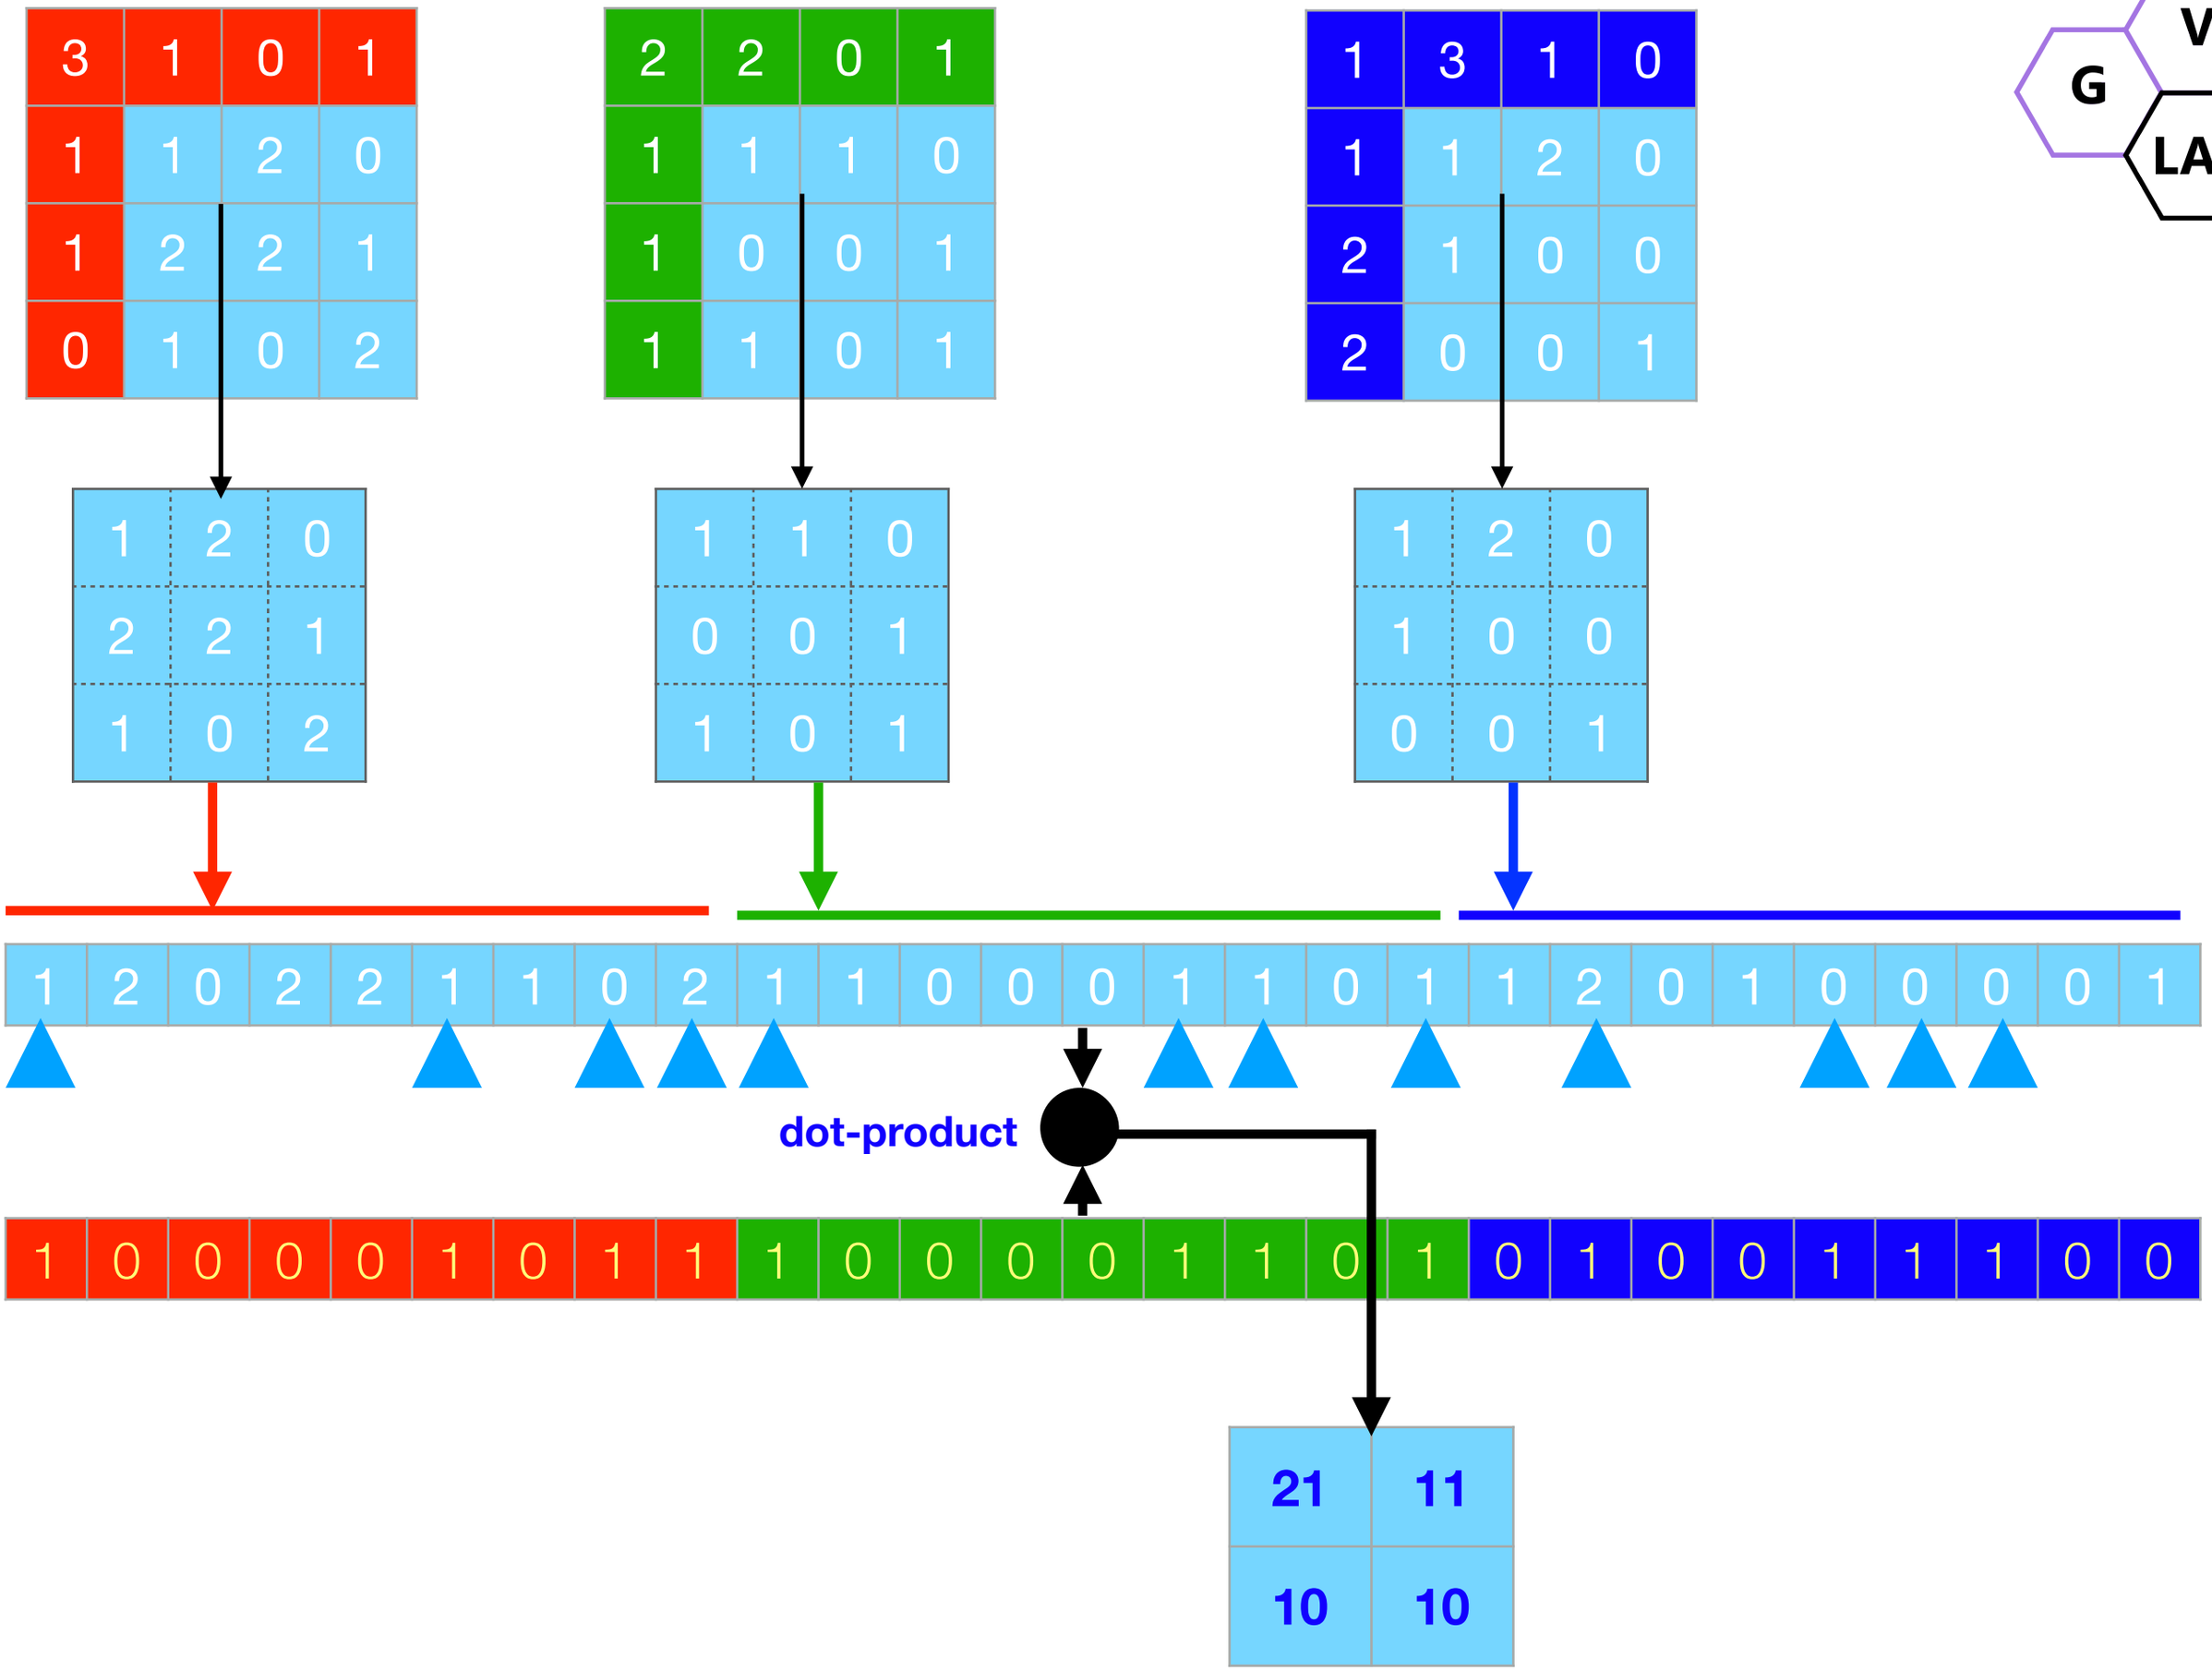
3 1 0 1
1 1 2 0
1 2 2 1
0 1 0 2
2 2 0 1
1 1 1 0
1 0 0 1
1 1 0 1
1 3 1 0
1 1 2 0
2 1 0 0
2 0 0 1
1 2 0
2 2 1
1 0 2
1 1 0
0 0 1
1 0 1
1 2 0
1 0 0
0 0 1
1 2 0 2 2 1 1 0 2 1 1 0 0 0 1 1 0 1 1 2 0 1 0 0 0 0 1
dot-product
1 0 0 0 0 1 0 1 1 1 0 0 0 0 1 1 0 1 0 1 0 0 1 1 1 0 0
21 11
10 10

Multi-channels input

Final result

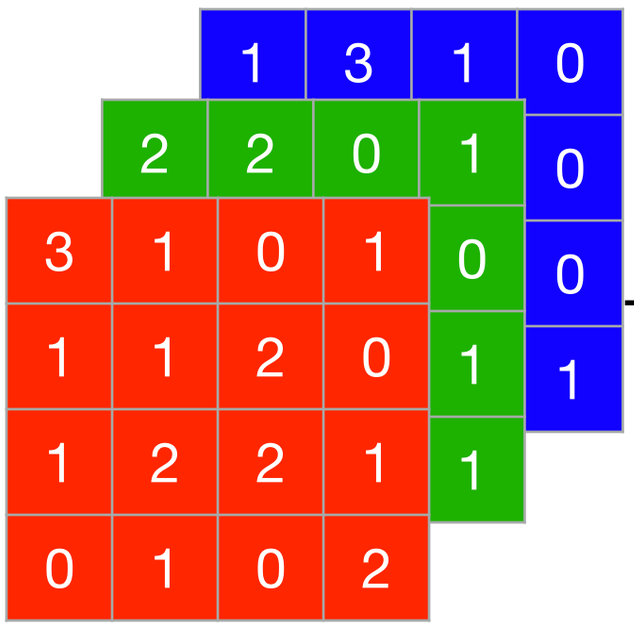

Input: 3 channels

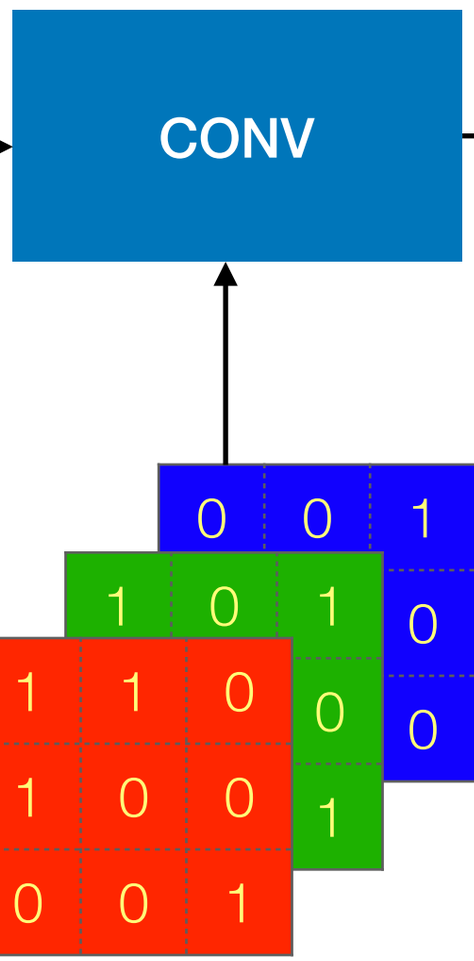

Kernel: 3 channels

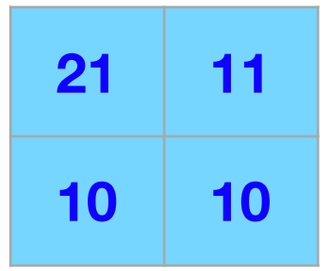

Output: 1 channel

Convolution

Multi-filters layer

Dr. Thanh-Sach LE
LTSACH@hcmut.edu.vn

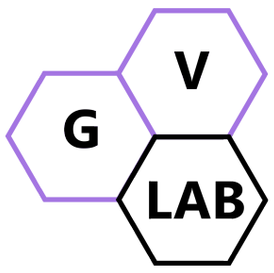

GVLab:
Graphics and Vision Laboratory

Faculty of Computer Science and Engineering,
HCMUT

20 Multi-filters layer

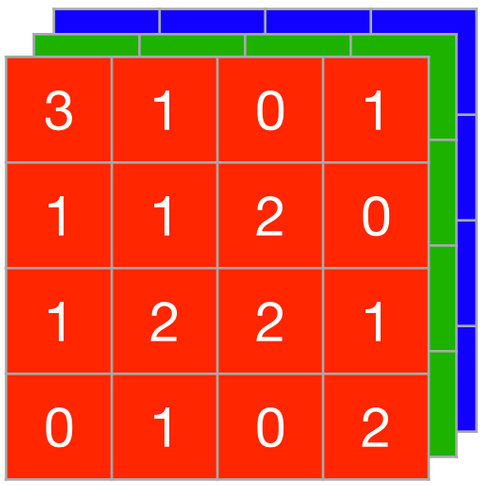

Input feature map
(D = 3 channels)

- Input: a feature map of D channels
- Convolution layer in Deep learning frameworks:
  - Consists of multiple filters:
    - Each the filters' kernel has D channels (#channels of the input)
    - All the filters' kernel are the same size, e.g., 3x3
- How the convolution layer computed?

21 Multi-filters layer

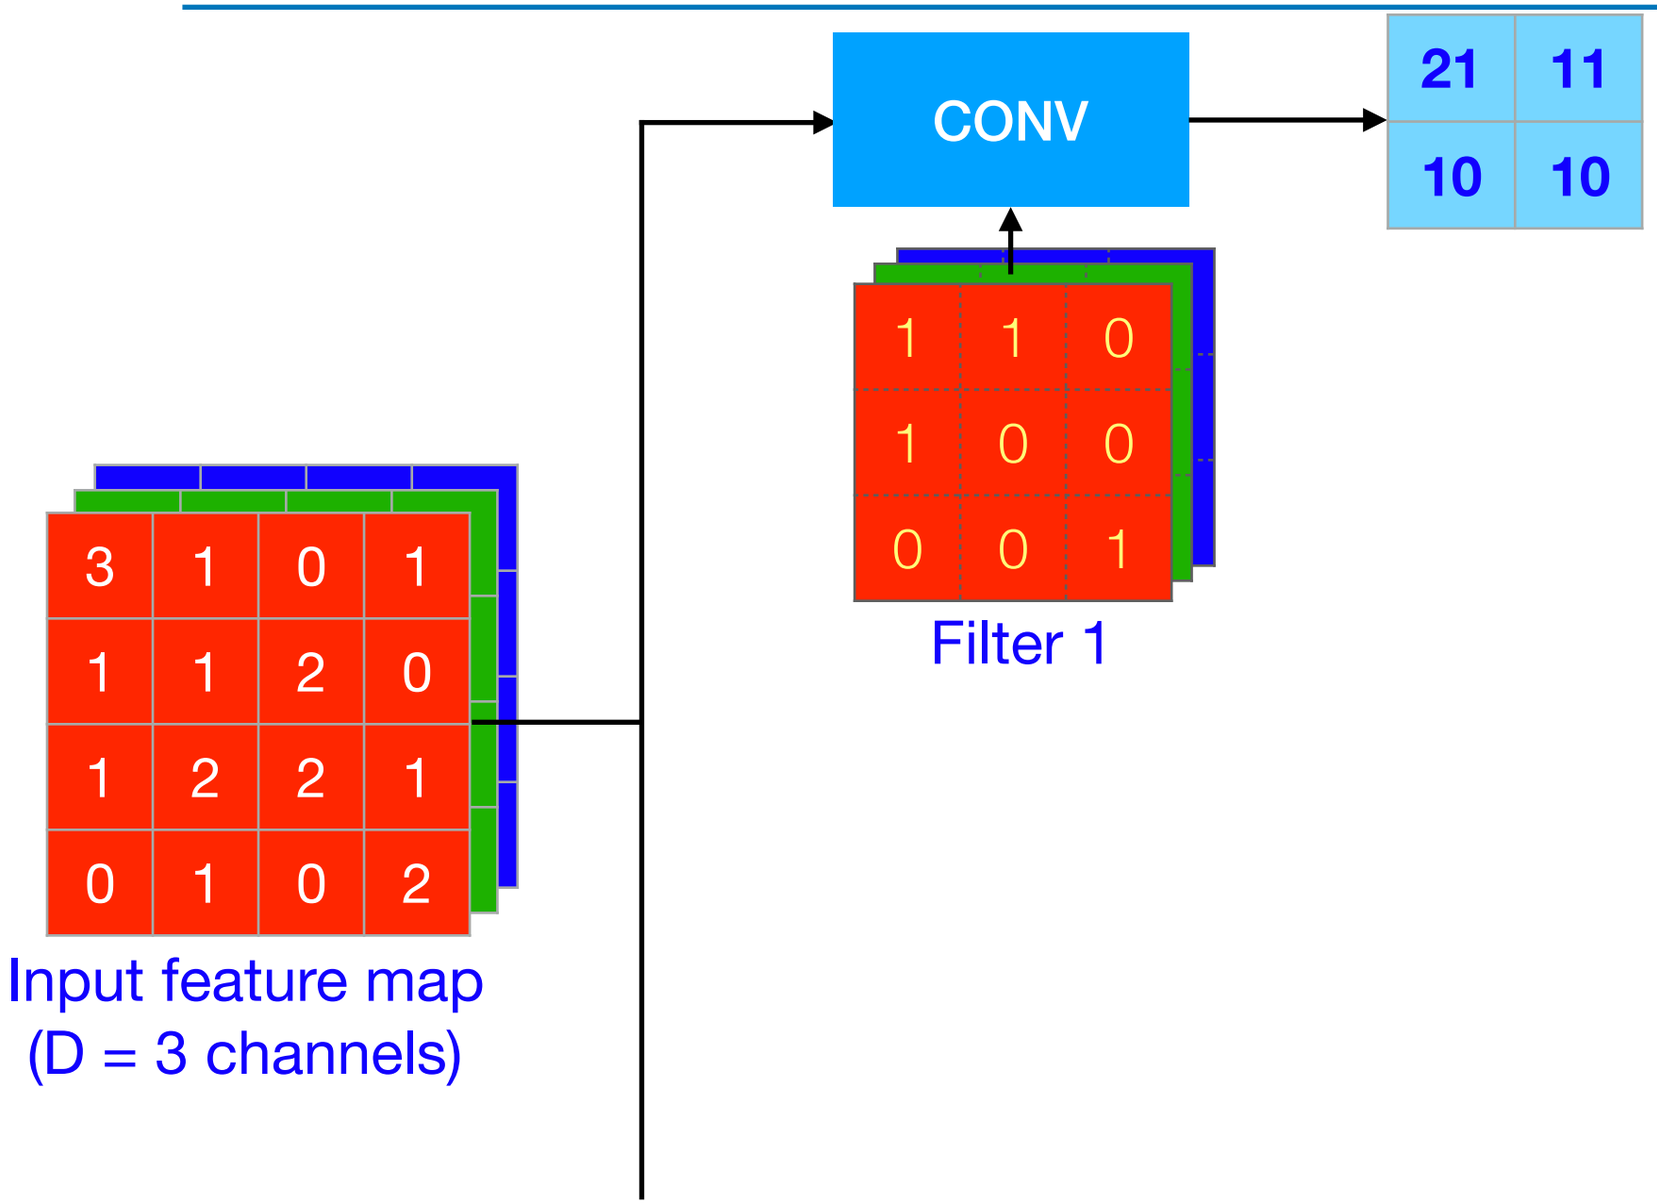

22 Multi-filters layer

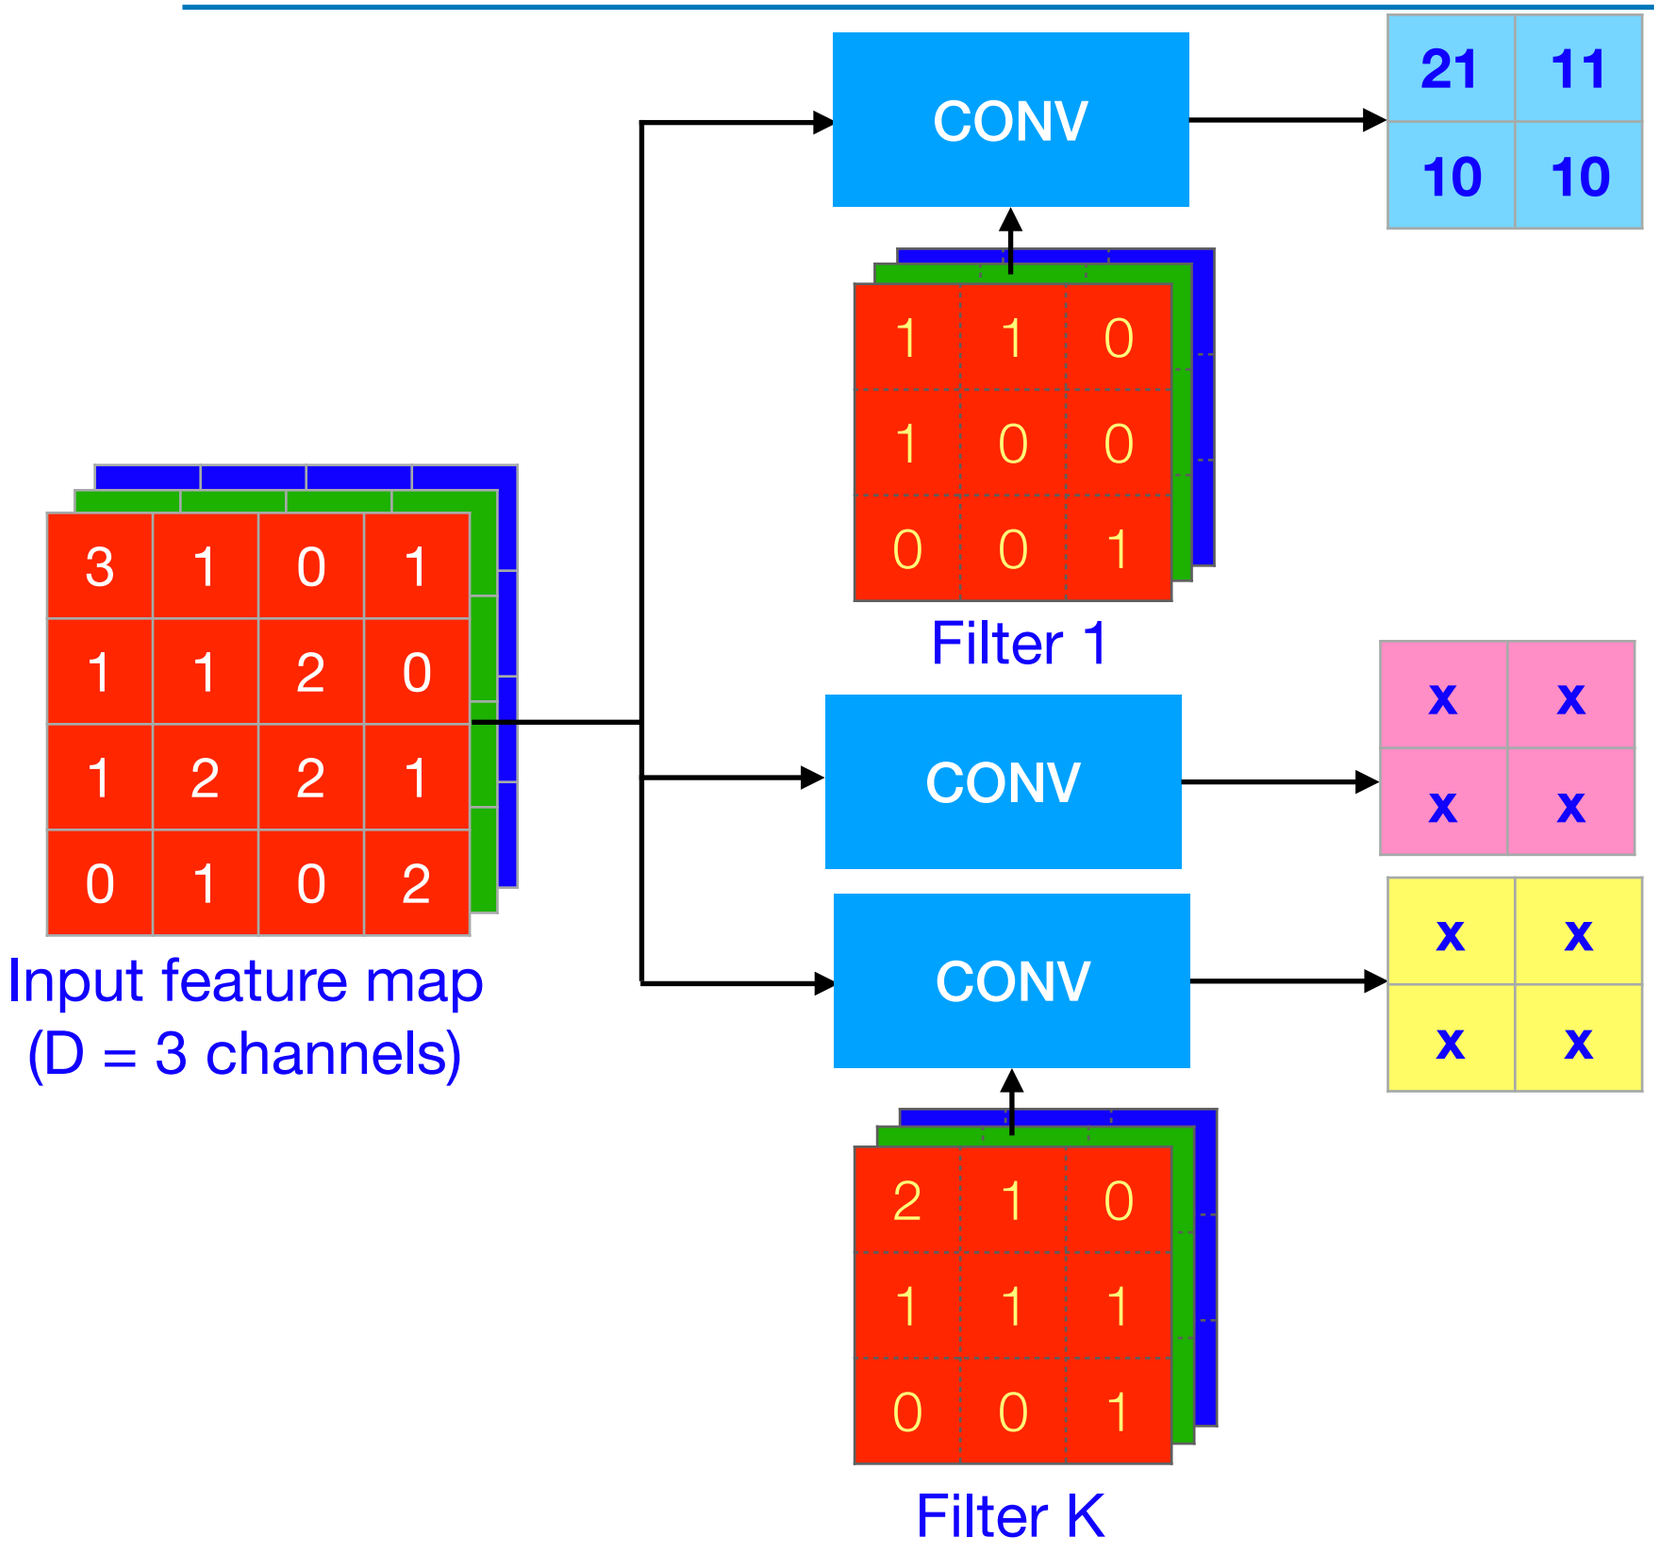

23 Multi-filters layer